

# Portfolio assignment 17

30 min: Train a decision tree to predict the body\_mass\_g of a penguin based on their characteristics.

- Split the penguin dataset into a train (70%) and test (30%) set.
- Use the train set to fit a DecisionTreeRegressor. You are free to to choose which columns you want to use as feature variables and you are also free to choose the max\_depth of the tree. **Note:** Some machine learning algorithms can not handle missing values. You will either need to
  - replace missing values (with the mean or most popular value). For replacing missing values you can use .fillna(<value>) <https://pandas.pydata.org/docs/reference/api/pandas.Series.fillna.html>
  - remove rows with missing data. You can remove rows with missing data with .dropna() <https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.dropna.html>
- Use your decision tree model to make predictions for both the train and test set.
- Calculate the RMSE for both the train set predictions and test set predictions.
- Is the RMSE different? Did you expect this difference?
- Use the plot\_tree\_regression function above to create a plot of the decision tree. Take a few minutes to analyse the decision tree. Do you understand the tree?

## Regression

When the target variable is a numerical variable then we refer to this task as a regression task.

Examples of regression tasks:

- Predict the price people are willing to pay for a house.
- Predict the salary a student will earn in the future.

Examples of machine learning algorithms that we could use for regression:

- Decision trees
- Random forests
- Linear regression
- Neural networks

In [1]:  
`import pandas as pd  
import seaborn as sns`

In [2]:  
`penguins = sns.load_dataset("penguins")  
penguins.head()`

Out[2]:

	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g	sex
0	Adelie	Torgersen	39.1	18.7	181.0	3750.0	Male
1	Adelie	Torgersen	39.5	17.4	186.0	3800.0	Female
2	Adelie	Torgersen	40.3	18.0	195.0	3250.0	Female
3	Adelie	Torgersen	NaN	NaN	NaN	NaN	NaN
4	Adelie	Torgersen	36.7	19.3	193.0	3450.0	Female

In [3]:  
`from sklearn.tree import DecisionTreeRegressor`

In [4]:  
`penguins = penguins.dropna() # removing a low amount of na's`

In [5]:  
`features= ['body_mass_g']  
dt_regression = DecisionTreeRegressor(max_depth = 3) # Increase max_depth to see effect  
dt_regression.fit(penguins[features], penguins['bill_length_mm'])`

Out[5]: DecisionTreeRegressor(max\_depth=3)

In [6]:  
`from sklearn import tree  
import graphviz  
  
def plot_tree_regression(model, features):  
 # Generate plot data  
 dot_data = tree.export_graphviz(model, out_file=None,  
 feature_names=features,  
 filled=True, rounded=True,  
 special_characters=True)  
  
 # Turn into graph using graphviz  
 graph = graphviz.Source(dot_data)  
  
 # Write out a pdf  
 graph.render("decision_tree")  
  
 # Display in the notebook  
 return graph`

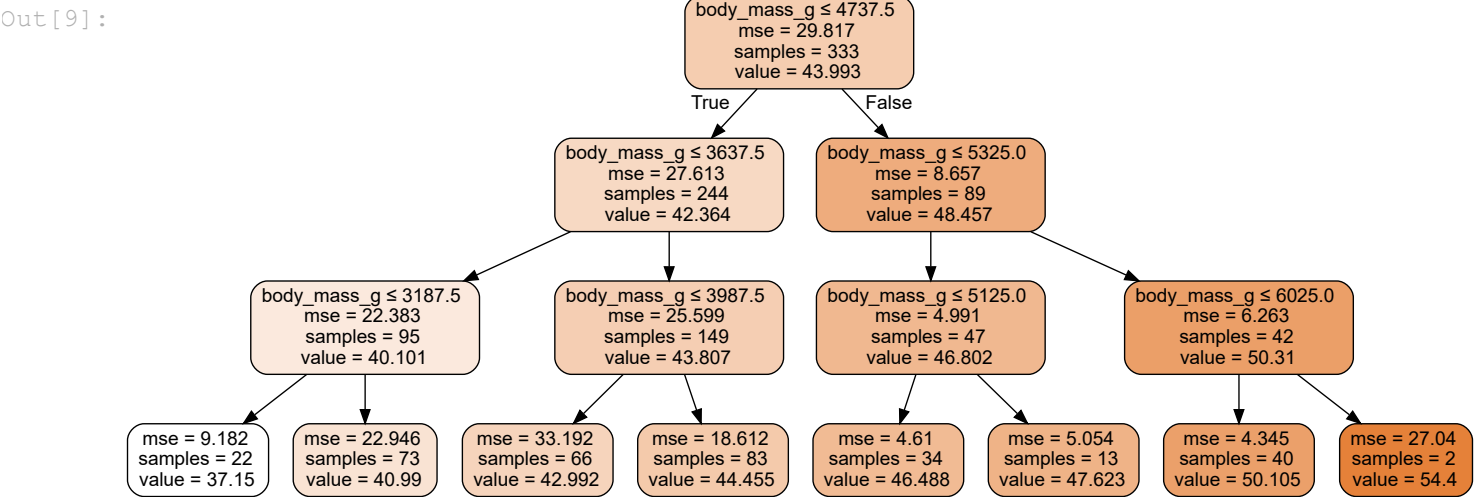
In [7]:  
`def calculate_rmse(predictions, actuals):  
 if(len(predictions) != len(actuals)):  
 raise Exception("The amount of predictions did not equal the amount of actuals")  
  
 return (((predictions - actuals) ** 2).sum() / len(actuals)) ** (1/2)`

In [8]:  
`predictionsOnTrainset = dt_regression.predict(penguins[features])  
predictionsOnTestset = dt_regression.predict(penguins[features])  
  
rmseTrain = calculate_rmse(predictionsOnTrainset, penguins.bill_length_mm)  
rmseTest = calculate_rmse(predictionsOnTestset, penguins.bill_length_mm)  
  
print("RMSE on training set " + str(rmseTrain))  
print("RMSE on test set " + str(rmseTest))`

RMSE on training set 4.266950627893799  
RMSE on test set 4.266950627893799

The difference seems to be very minimal, this means the training set can probably predict the output very well. The errors SHOULD be minimal, but you're never sure.

In [9]:  
`plot_tree_regression(dt_regression, features)`



In [ ]: